

CNN Image Classification

```
import numpy as np
import tensorflow as tf
from tensorflow import keras
from keras.models import Sequential
import pandas as pd
from matplotlib import pyplot as plt
%matplotlib inline
import os
import cv2
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, GlobalAveragePooling2D
from PIL import Image
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix, classification_report
```

```
# Define the paths to your image and csv folders
train_val_dir = "C:\\Users\\335224\\Downloads\\Chart\\train_val"
test_dir = "C:\\Users\\335224\\Downloads\\Chart\\test"
train_path_labels = "C:\\Users\\335224\\Downloads\\Chart\\train_v"
train_val_labels = pd.read_csv(train_path_labels)
```

```
# Load training dataset in numpy array

images = []
labels = []

for filename in os.listdir(train_val_dir):
    if filename.endswith('.png'):
        # Load the images and resize them to (128, 128) with 3 color channels
        img = cv2.imread(os.path.join(train_val_dir, filename))
        img = cv2.resize(img, (128, 128))
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

#         img = Image.open(os.path.join(train_val_dir, filename))
#         img_array = np.array(img)
#         # Append the array to the list of images
#         images.append(img_array)
#         labels.append(filename)

# Convert the string labels to numerical labels
le = LabelEncoder()
labels = le.fit_transform(labels)
```

```
# Convert the lists to NumPy arrays
images = np.array(images)
labels = np.array(labels)
# Save the arrays in NumPy format
np.save('x_train.npy', images)
np.save('y_train.npy', labels)
x_train = np.load('x_train.npy')
y_train = np.load('y_train.npy')
```

```
x_train.shape
```

```
(1000, 128, 128, 3)
```

```
x_train[:5]
y_train[:5]
```

```
array([0, 1, 2, 3, 4], dtype=int64)
```

```
# Load test dataset in numpy array

images = []
labels = []

for filename in os.listdir(test_dir):
    if filename.endswith('.png'):
        # Load the images and resize them to (128, 128) with 3 color channels
        img = cv2.imread(os.path.join(test_dir, filename))
        img = cv2.resize(img, (128, 128))
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

#     img = Image.open(os.path.join(test_dir, filename))
#     img_array = np.array(img)
#     # Append the array to the list of images
#     images.append(img_array)
#     labels.append(filename)

# Convert the string labels to numerical labels
le = LabelEncoder()
labels = le.fit_transform(labels)

# Convert the lists to NumPy arrays
```

```

images = np.array(images)
labels = np.array(labels)

# Save the arrays in NumPy format
np.save('x_test.npy', images)
np.save('y_test.npy', labels)

x_test = np.load('x_test.npy')
y_test = np.load('y_test.npy')

```

```
x_test.shape
```


```
(50, 128, 128, 3)
```

```

# check the images loaded
plt.figure(figsize = (10,2))
plt.imshow(x_train[10])
plt.imshow(x_train[208])
plt.imshow(x_train[444])

```

```
<matplotlib.image.AxesImage at 0x1b79ae3b730>
```



```

# define some classes from the images we have observed
image_classes = ['line', 'dot_line', 'hbar_categorical', 'vbar_categorical', 'pie']
image_classes[0]

# map the categories to the labels array i.e y_train
label_map = {'line': 0, 'dot_line': 1, 'hbar_categorical': 2, 'vbar_categorical': 3, 'pie': 4}
y_train = np.array([label_map[label] for label in train_val_labels['type']])

```


```
y_train
y_train.shape
y_test.shape
```


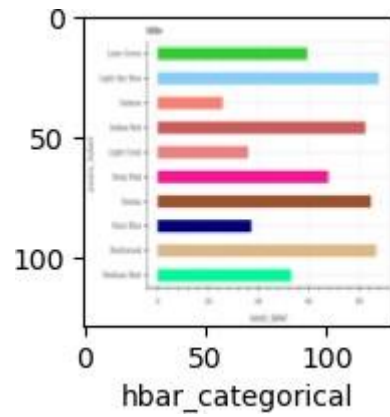
(50,)

```
# function to test the chart sample
```

```
def image_sample(x, y, index):
    plt.figure(figsize = (10,2))
    plt.imshow(x[index])
    #image_label = train_val_labels.iloc[index]['type']
    #plt.xlabel(image_label)
    plt.xlabel(image_classes[y[index]])
```

```
image_sample(x_train,y_train,0)
image_sample(x_train,y_train,208)
image_sample(x_train,y_train,444)
```





```
# normalize the image
```

```
x_train=x_train /255
```

```
x_test=x_train /255
```

```
x_test.shape
```

```
(1000, 128, 128, 3)
```

```
# take the label for train data from csv file
```

```
y_train_index = train_val_labels['image_index']  
y_train_type = train_val_labels['type']
```

```
y_train_type[:5]
```

```
0    vbar_categorical  
1    vbar_categorical  
2    vbar_categorical  
3    vbar_categorical  
4    vbar_categorical  
Name: type, dtype: object
```

```
# Define the model architecture
```

```
model = Sequential([  
    Flatten(input_shape=(128,128,3)),  
    Dense(3000, activation='relu'),  
    Dense(1000, activation='relu'),  
    Dense(5, activation='softmax')  
])
```

```
# Compile the model
```

```
model.compile(optimizer='SGD', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
model.fit(x_train,y_train,epochs=10)
```

```

Epoch 1/10
32/32 [=====] - 16s 473ms/step - loss: 7.7490 - accuracy: 0.1940
Epoch 2/10
32/32 [=====] - 15s 475ms/step - loss: 1.6554 - accuracy: 0.2430
Epoch 3/10
32/32 [=====] - 15s 474ms/step - loss: 1.6163 - accuracy: 0.2260
Epoch 4/10
32/32 [=====] - 15s 469ms/step - loss: 1.6223 - accuracy: 0.2020
Epoch 5/10
32/32 [=====] - 15s 473ms/step - loss: 1.6105 - accuracy: 0.2300
Epoch 6/10
32/32 [=====] - 15s 471ms/step - loss: 1.6102 - accuracy: 0.2120
Epoch 7/10
32/32 [=====] - 15s 472ms/step - loss: 1.6109 - accuracy: 0.2030
Epoch 8/10
32/32 [=====] - 15s 472ms/step - loss: 1.6130 - accuracy: 0.1970
Epoch 9/10
32/32 [=====] - 15s 473ms/step - loss: 1.6102 - accuracy: 0.1960
Epoch 10/10
32/32 [=====] - 15s 472ms/step - loss: 1.6099 - accuracy: 0.2030

```

<keras.callbacks.History at 0x1b79a6714c0>

```

# Split the training images and labels into training and validation sets
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x_train, y_train, test_size=0.2, random_state=42)

```

```
model.evaluate(x_test,y_test)
```

```

7/7 [=====] - 1s 106ms/step - loss: 1.6102 - accuracy: 0.1650
[1.6102104187011719, 0.16500000655651093]

```

```

y_pred = model.predict(x_test)
y_pred
y_pred_classes = [np.argmax(ele) for ele in y_pred]
# print("classification report : \n",classification_report(y_test,y_pred_classes))

```

```
7/7 [=====] - 1s 111ms/step
```

here we see the accuracy is very low and we need to modify our nn to add more layers for better accuracy

```
# Print the shapes of the arrays to verify that they are loaded correctly
print("Train Images Shape:", x_train.shape)
print("Train Labels Shape:", y_train.shape)
print("Test Images Shape:", x_test.shape)
print("Test Labels Shape:", y_test.shape)
```

```
Train Images Shape: (800, 128, 128, 3)
Train Labels Shape: (800,)
Test Images Shape: (200, 128, 128, 3)
Test Labels Shape: (200,)
```

```
# modifying the model architecture to cnn
cnn_model = Sequential([
    Conv2D(filters=16, kernel_size=(3,3), activation='relu', input_shape=(128,128,3)),
    MaxPooling2D(pool_size=(2,2)),
    Conv2D(32, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(5, activation='softmax')
])

cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = cnn_model.fit(x_train, y_train, batch_size=1000, epochs=50, validation_data=(x_test, y_test))


# Plot the obtained loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.show()
```



```
1/1 [=====] - 14s 14s/step - loss: 1.6346 - accuracy: 0.1950 - val_loss: 2.6084 - val_accuracy: 0.1650
Epoch 2/50
1/1 [=====] - 6s 6s/step - loss: 2.3806 - accuracy: 0.2087 - val_loss: 1.7656 - val_accuracy: 0.2950
Epoch 3/50
1/1 [=====] - 6s 6s/step - loss: 1.7012 - accuracy: 0.3050 - val_loss: 1.6191 - val_accuracy: 0.1900
Epoch 4/50
1/1 [=====] - 6s 6s/step - loss: 1.6281 - accuracy: 0.2025 - val_loss: 1.5837 - val_accuracy: 0.3500
Epoch 5/50
1/1 [=====] - 6s 6s/step - loss: 1.5960 - accuracy: 0.3050 - val_loss: 1.5716 - val_accuracy: 0.3250
Epoch 6/50
1/1 [=====] - 7s 7s/step - loss: 1.5638 - accuracy: 0.3013 - val_loss: 1.5697 - val_accuracy: 0.3000
Epoch 7/50
1/1 [=====] - 6s 6s/step - loss: 1.5404 - accuracy: 0.3600 - val_loss: 1.5178 - val_accuracy: 0.4600
Epoch 8/50
1/1 [=====] - 7s 7s/step - loss: 1.4966 - accuracy: 0.5200 - val_loss: 1.5030 - val_accuracy: 0.3500
Epoch 9/50
1/1 [=====] - 9s 9s/step - loss: 1.4610 - accuracy: 0.3862 - val_loss: 1.4351 - val_accuracy: 0.4850
Epoch 10/50
1/1 [=====] - 6s 6s/step - loss: 1.4076 - accuracy: 0.5425 - val_loss: 1.4130 - val_accuracy: 0.4400
Epoch 11/50
1/1 [=====] - 6s 6s/step - loss: 1.3648 - accuracy: 0.4900 - val_loss: 1.3407 - val_accuracy: 0.5250
Epoch 12/50
1/1 [=====] - 6s 6s/step - loss: 1.3073 - accuracy: 0.5575 - val_loss: 1.2953 - val_accuracy: 0.5250
Epoch 13/50
1/1 [=====] - 6s 6s/step - loss: 1.2554 - accuracy: 0.5600 - val_loss: 1.2201 - val_accuracy: 0.5550
Epoch 14/50
1/1 [=====] - 6s 6s/step - loss: 1.2049 - accuracy: 0.5412 - val_loss: 1.1901 - val_accuracy: 0.5450
Epoch 15/50
1/1 [=====] - 6s 6s/step - loss: 1.1631 - accuracy: 0.5387 - val_loss: 1.1579 - val_accuracy: 0.5600
Epoch 16/50
1/1 [=====] - 6s 6s/step - loss: 1.1340 - accuracy: 0.5325 - val_loss: 1.2010 - val_accuracy: 0.5400
Epoch 17/50
1/1 [=====] - 6s 6s/step - loss: 1.1389 - accuracy: 0.5250 - val_loss: 1.2159 - val_accuracy: 0.5200
Epoch 18/50
1/1 [=====] - 6s 6s/step - loss: 1.2253 - accuracy: 0.4750 - val_loss: 1.1012 - val_accuracy: 0.5850
Epoch 19/50
1/1 [=====] - 6s 6s/step - loss: 1.0679 - accuracy: 0.5575 - val_loss: 1.2918 - val_accuracy: 0.4900
Epoch 20/50
1/1 [=====] - 6s 6s/step - loss: 1.1772 - accuracy: 0.5200 - val_loss: 1.0819 - val_accuracy: 0.6400
Epoch 21/50
1/1 [=====] - 6s 6s/step - loss: 1.0278 - accuracy: 0.6300 - val_loss: 1.1356 - val_accuracy: 0.5150
Epoch 22/50
1/1 [=====] - 6s 6s/step - loss: 1.0737 - accuracy: 0.5550 - val_loss: 1.1331 - val_accuracy: 0.5200
```

1/1 [=====]	- 6s	6s/step	- loss: 1.0418	- accuracy: 0.6100	- val_loss: 1.1245	- val_accuracy: 0.5450
Epoch 24/50						
1/1 [=====]	- 6s	6s/step	- loss: 1.0024	- accuracy: 0.5800	- val_loss: 1.1334	- val_accuracy: 0.5600
Epoch 25/50						
1/1 [=====]	- 6s	6s/step	- loss: 1.0057	- accuracy: 0.5913	- val_loss: 1.1035	- val_accuracy: 0.5950
Epoch 26/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.9845	- accuracy: 0.5938	- val_loss: 1.0920	- val_accuracy: 0.6050
Epoch 27/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.9554	- accuracy: 0.6400	- val_loss: 1.1239	- val_accuracy: 0.5650
Epoch 28/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.9514	- accuracy: 0.6550	- val_loss: 1.1200	- val_accuracy: 0.5250
Epoch 29/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.9308	- accuracy: 0.6313	- val_loss: 1.0869	- val_accuracy: 0.5500
Epoch 30/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.9033	- accuracy: 0.6263	- val_loss: 1.0772	- val_accuracy: 0.5750
Epoch 31/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.8984	- accuracy: 0.6500	- val_loss: 1.0723	- val_accuracy: 0.6050
Epoch 32/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.8794	- accuracy: 0.6725	- val_loss: 1.0841	- val_accuracy: 0.5700
Epoch 33/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.8572	- accuracy: 0.6875	- val_loss: 1.1049	- val_accuracy: 0.5650
Epoch 34/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.8499	- accuracy: 0.6825	- val_loss: 1.0810	- val_accuracy: 0.5850
Epoch 35/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.8264	- accuracy: 0.6825	- val_loss: 1.0673	- val_accuracy: 0.5950
Epoch 36/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.8075	- accuracy: 0.6975	- val_loss: 1.0948	- val_accuracy: 0.6050
Epoch 37/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.7960	- accuracy: 0.6963	- val_loss: 1.1023	- val_accuracy: 0.5850
Epoch 38/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.7769	- accuracy: 0.6988	- val_loss: 1.0961	- val_accuracy: 0.6000
Epoch 39/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.7621	- accuracy: 0.7225	- val_loss: 1.0868	- val_accuracy: 0.6050
Epoch 40/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.7451	- accuracy: 0.7250	- val_loss: 1.1040	- val_accuracy: 0.6100
Epoch 41/50						
1/1 [=====]	- 5s	5s/step	- loss: 0.7294	- accuracy: 0.7225	- val_loss: 1.1236	- val_accuracy: 0.6050
Epoch 42/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.7186	- accuracy: 0.7350	- val_loss: 1.1163	- val_accuracy: 0.5900
Epoch 43/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.7050	- accuracy: 0.7287	- val_loss: 1.1502	- val_accuracy: 0.5700
Epoch 44/50						
1/1 [=====]	- 6s	6s/step	- loss: 0.6954	- accuracy: 0.7387	- val_loss: 1.1710	- val_accuracy: 0.5900

Epoch 45/50
1/1 [=====] - 5s 5s/step - loss: 0.6979 - accuracy: 0.7163 - val_loss: 1.2036 - val_accuracy: 0.5450
Epoch 46/50
1/1 [=====] - 6s 6s/step - loss: 0.7217 - accuracy: 0.7063 - val_loss: 1.1758 - val_accuracy: 0.5900
Epoch 47/50
1/1 [=====] - 6s 6s/step - loss: 0.6455 - accuracy: 0.7575 - val_loss: 1.2177 - val_accuracy: 0.5450
Epoch 48/50
1/1 [=====] - 5s 5s/step - loss: 0.6721 - accuracy: 0.7300 - val_loss: 1.1955 - val_accuracy: 0.5650
Epoch 49/50
1/1 [=====] - 5s 5s/step - loss: 0.6645 - accuracy: 0.7375 - val_loss: 1.1970 - val_accuracy: 0.5850
Epoch 50/50
1/1 [=====] - 6s 6s/step - loss: 0.6303 - accuracy: 0.7650 - val_loss: 1.2362 - val_accuracy: 0.5850

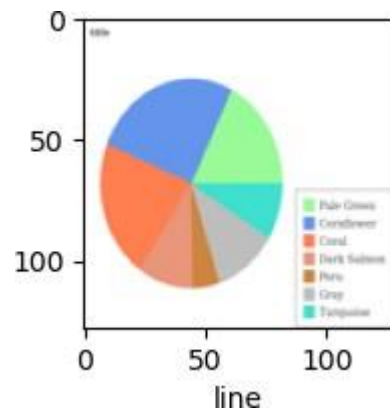



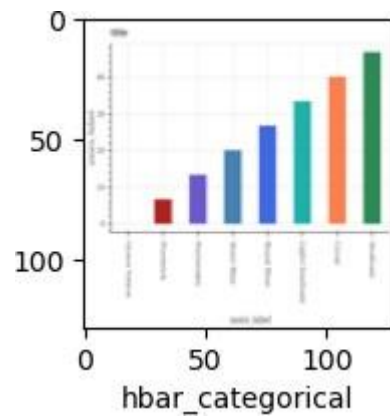

```
cnn_model.evaluate(x_test,y_test)
```

```
7/7 [=====] - 0s 41ms/step - loss: 1.2362 - accuracy: 0.5850
```

[1.2362478971481323, 0.5849999785423279]

```
image_sample(x_test,y_test,1)  
image_sample(x_test,y_test,50)  
image_sample(x_test,y_test,25)  
image_sample(x_test,y_test,30)
```





Observation: we are able to see some wrong predictions here

```
y_pred = cnn_model.predict(x_test)
y_pred[:5]
```

7/7 [=====] - 0s 39ms/step

```
array([[1.88634619e-01, 4.53442723e-01, 1.89865723e-01, 3.13057527e-02,
        1.36751115e-01],
       [1.16687104e-01, 1.73624381e-02, 6.09002418e-05, 7.50791188e-03,
        8.58381629e-01],
       [4.57645766e-02, 1.14214392e-02, 9.25663917e-05, 2.46962626e-03,
        9.40251827e-01],
       [5.41297436e-01, 5.38021093e-04, 9.51929204e-03, 3.77870835e-02,
        4.10858095e-01],
       [3.89628053e-01, 3.40805709e-01, 1.78500041e-01, 1.64468344e-02,
        7.46192932e-02]], dtype=float32)
```

```
y_classes = [np.argmax(element) for element in y_pred]
y_classes[:5]
```

```
[1, 4, 4, 0, 0]
```

```
y_test[:5]
```

```
array([0, 4, 4, 4, 0])
```

```
# we can see some values are not matching here
```

```
# test actual and predicted
```

```
# image_sample(x_test,y_test,1) #actual
# image_classes[y_classes[1]] #predicted
```

```
# image_sample(x_test,y_test,10) #actual
# image_classes[y_classes[10]] #predicted
```

```
image_sample(x_test,y_test,15) #actual
image_classes[y_classes[15]] #predicted
```


```
'dot_line'
```


Confusion Matrix:

```
[[ 9 17  0  4  7]
 [ 3 35  0  3  3]
 [ 5  5 21 17  0]
 [ 0  3  5 23  2]
 [ 4  1  0  4 29]]
```

```
# Plot the confusion matrix
import seaborn as sn
plt.figure(figsize = (10,10))
sn.heatmap(conf_mat,annot=True,fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
Text(95.7222222222221, 0.5, 'Actual')
```



```
# for 50 iterations, we can see some promising accuracy, more training will be required for better accuracy  
# in the confusion matrix, whatever is not in diagonal is considered an error
```

```
from tensorflow.keras.applications import VGG16  
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
# Load the pre-trained model  
vgg16_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
```

```
# Replace the final classification layer with a new layer  
x = vgg16_model.output  
x = GlobalAveragePooling2D()(x)  
x = Dense(128, activation='relu')(x)  
  
predictions = Dense(5, activation='softmax')(x)  
pt_model = tf.keras.Model(inputs=vgg16_model.input, outputs=predictions)
```

```
# Freeze the weights of all layers except the new classification layer  
for layer in pt_model.layers:  
    layer.trainable = False
```

```
# Compile the model with categorical crossentropy loss and Adam optimizer  
pt_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
# Print the summary of the model architecture  
pt_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0

dense_5 (Dense)	(None, 128)	65664
dense_6 (Dense)	(None, 5)	645

=====
Total params: 14,780,997
Trainable params: 0
Non-trainable params: 14,780,997

Set up data generators for image augmentation and feeding data to the model

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=20,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')  
  
test_datagen = ImageDataGenerator(rescale=1./255)
```

```
train_generator = train_datagen.flow(x_train, y_train, batch_size=32)  
  
test_generator = train_datagen.flow(x_test, y_test, batch_size=32)
```

Training the model with early stopping

```
from tensorflow.keras.callbacks import EarlyStopping  
  
es = EarlyStopping(monitor='val_loss', patience=10, verbose=1, mode='min', restore_best_weights=True)  
  
history = pt_model.fit(train_generator, epochs=100, validation_data=test_generator, callbacks=[es])
```

Epoch 1/100